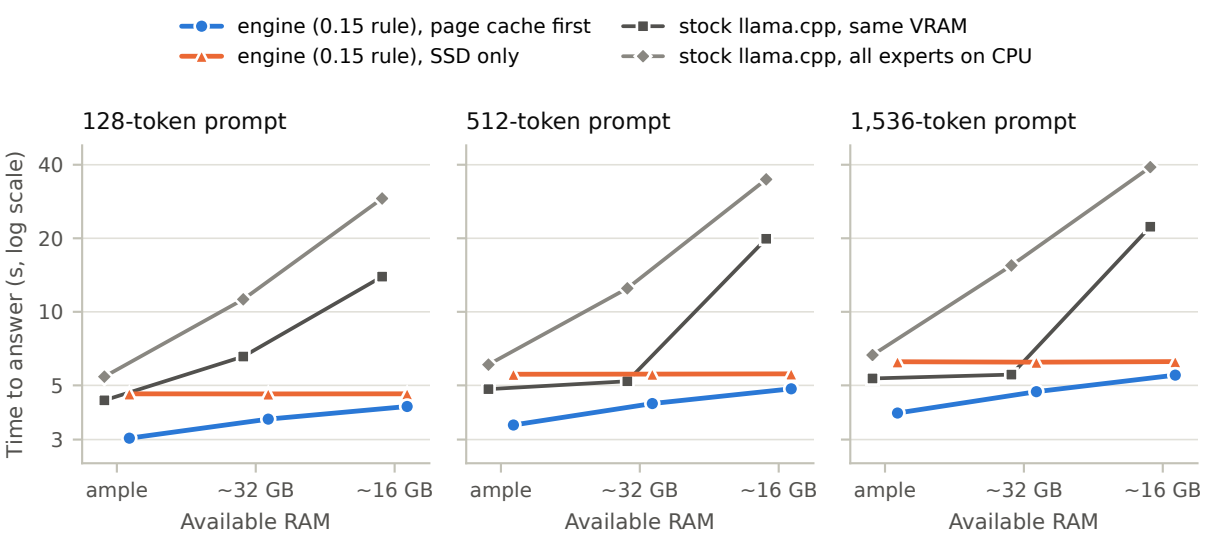
engine (0.15 rule), page cache first
engine (0.15 rule), SSD only
stock llama.cpp, same VRAM
stock llama.cpp, all experts on CPU
128-token prompt
512-token prompt
1,536-token prompt
Time to answer (s, log scale)
40
20
10
5
3
ample
~32 GB
~16 GB
Available RAM